

# Operating Systems

## 13. File Systems

Paul Krzyzanowski

Rutgers University

Spring 2015

# Terminology

# What's a file system?

- Traditionally
  - A way to manage variable-size persistent data
    - Organize, store, retrieve, delete information
  - Random access
    - Arbitrary files can be accessed by name
    - Arbitrary parts of a file can be accessed
  - File systems are implemented on top of block devices
- More abstract
  - A way to access information by name
    - Devices
    - System configuration, process info, random numbers

# Terms

- **Disk**
  - Non-volatile block-addressable storage.
- **Block = sector**
  - Smallest chunk of I/O on a disk
  - Common block sizes = 512 or 4096 (4K) bytes  
E.g., WD Black Series 4TB drive has 7,814,037,168 512-byte sectors
- **Partition**
  - Set of contiguous blocks on a disk. A disk has  $\geq 1$  partitions
- **Volume**
  - Disk, disks, or partition that contains a file system
  - A volume may span disks

# More terms

---

- **Track**
  - Blocks are stored on concentric *tracks* on a disk
- **Cylinder**
  - The set of all blocks on one track  
(obsolete now since we don't know what's where)
- **Seek**
  - The movement of a disk head from track to track

# File Terms

---

- **File**
  - A unit of data managed by the file system
- **Data: (Contents)**
  - 
  - Unstructured (byte stream) or structured (records)
- **Name**
  - A textual name that identifies the file

# File Terms

---

- **Metadata**
  - Information about the file (creation time, permissions, length of file data, location of file data, etc.)
- **Attribute**
  - A form of metadata – a textual name and associated value (e.g., source URL, author of document, checksum)
- **Directory (folder)**
  - A container for file names
  - Directories within directories provide a hierarchical name system

# File System Terms

- **Superblock**
  - Area on the volume that contains key file system information
- **inode (file control block)**
  - A structure that stores a file's metadata and location of file data
- **Cluster**
  - Logical block size used in the file system that is equivalent to  $N$  blocks
- **Extent**
  - Group of contiguous clusters identified by a starting block number and a block count



# Design Choices

## Namespace

Flat, hierarchical, or other?

## Multiple volumes

Explicit device identification  
(A:, B:, C:, D:)

or integrate into one namespace?

## File types

Unstructured  
(byte streams)

or structured  
(e.g., indexed files)?

## File system types

Support one type of file system

or multiple types  
(iso9660, NTFS, ext3)?

## Metadata

What kind of attributes should the file system have?

## Implementation

How is the data laid out on the disk?

# Working with the Operating System

## File System Operations

# Mounting

- Make file system available for use
- *mount* system call
  - Pass the file system type, block device & mount point
- Steps
  - Access the raw disk (block device)
  - Read superblock and file system metadata (free block bitmaps, root directory, etc.)
  - Check to see if the file system was properly unmounted (clean?)
    - If not, validate the structure of the file system
  - Prepare in-memory data structures to access the volume
    - In-memory version of the superblock
    - References to the root directory
    - Free block bitmaps
  - Mark the superblock as “dirty”

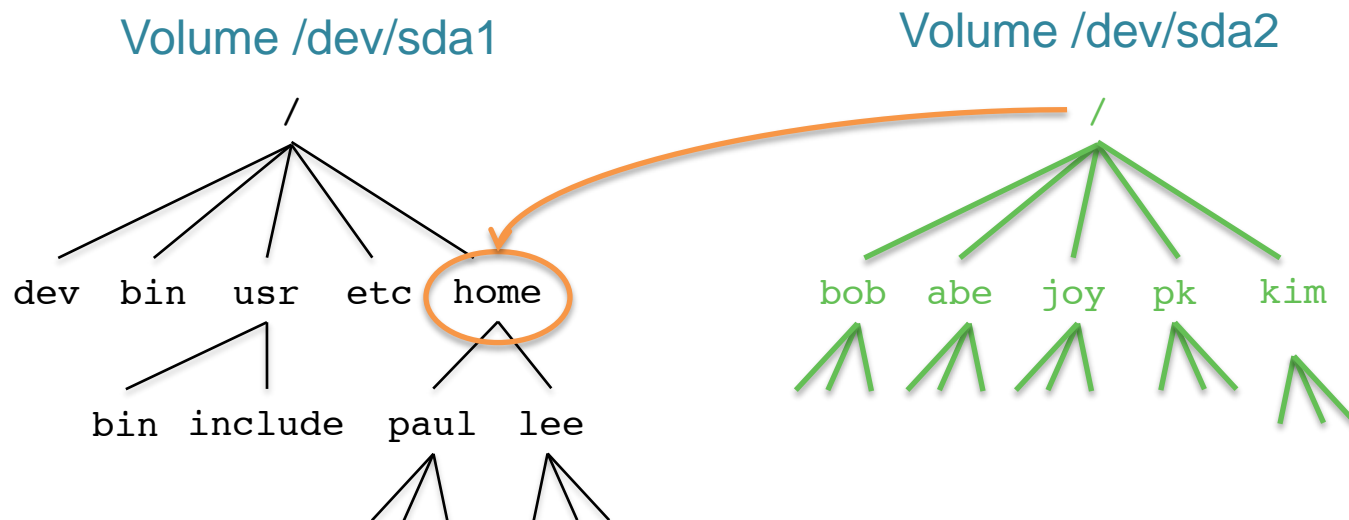
# Unmounting

---

- Ensure there are no processes with open files in the file system
- Remove file system from the OS name space
- Flush all in-memory file system state to disk
- Mark the superblock as “clean” (unmount took place)

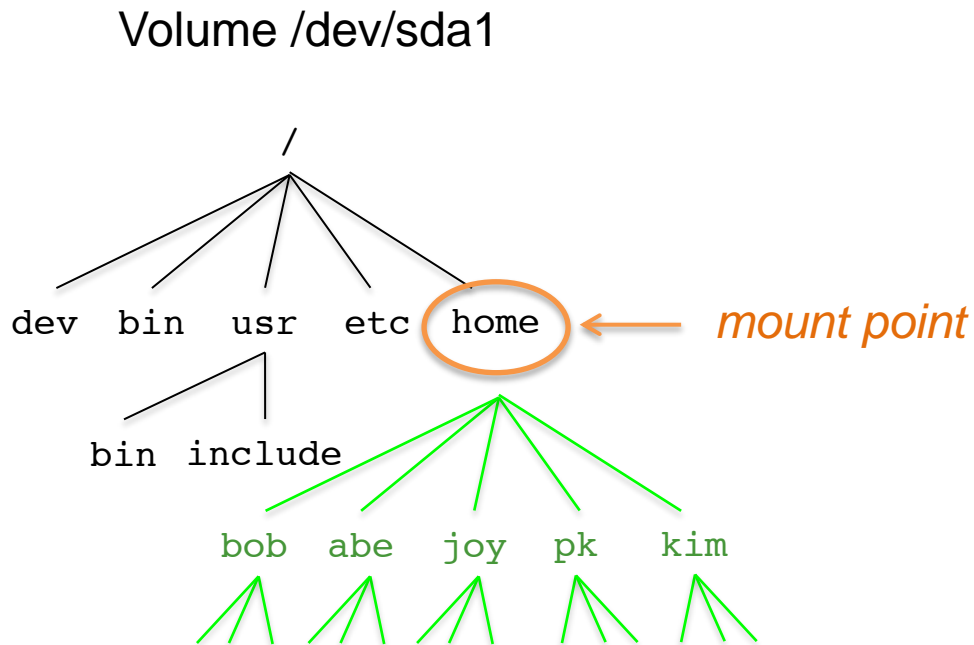
# Mounting: building up a name space

- Combine multiple file systems into a single hierarchical name space
- The mounted file system overlays (& hides) anything in the file system under that **mount point**
- Looking up a **pathname** may involve traversing multiple mount points



```
mount -t ext4 /dev/sda2 /home
/home becomes a mount point for /dev/sda2
```


# Mounting: building up a name space



`/home/paul` and `/home/lee` are no longer visible

# Union mounts

Mounted file system merges the existing namespace



Considerations:

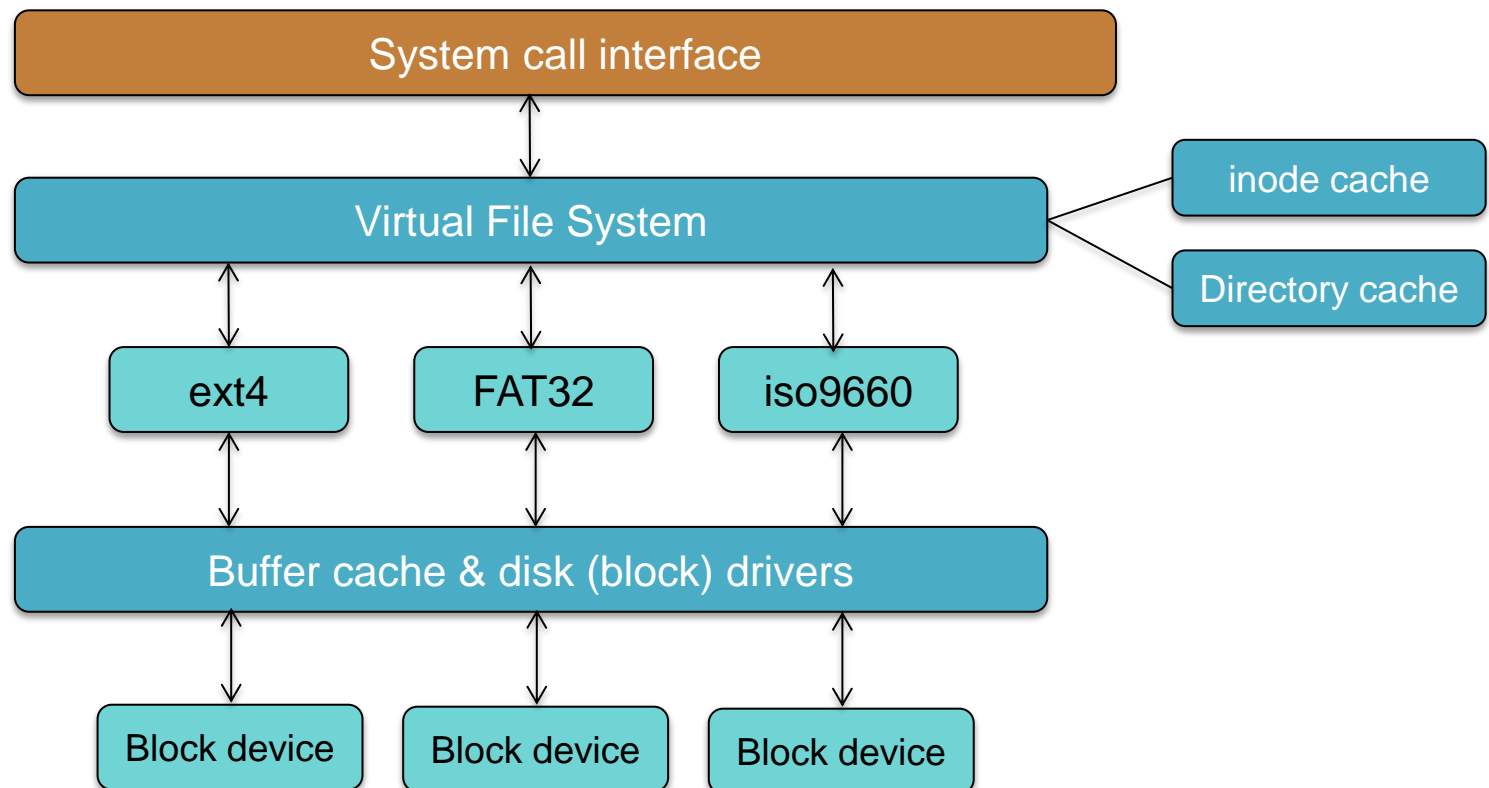
- Search path (what if two names are the same in the file systems)?
- Where to write?

# Operating System Interfaces for File Systems



# Virtual File System (VFS) Interface

- Abstract interface for a file system object
- Each *real* file system interface exports a common interface



# VFS: Common set of objects

- **Superblock**: Describes the file system
  - Block size, max file size, mount point
  - One per mounted file system
- **inode**: represents a single file
  - Unique identifier for every object (file) in a specific file system
  - File systems have methods to translate a name to an inode
  - VFS inode defines all the operations possible on it
- **dentry**: directory entries & contents
  - Name of file/directory, child dentries, parent
  - Directory entries: translations of names to inodes
- **file**: represents an open file
  - VFS keeps state: mode, read/write offset, etc.

# VFS superblock

- Structure that represents info about the file system
- Includes
  - File system name
  - Size
  - State
  - Reference to the block device
  - List of operations for managing inodes within the file system:
    - *alloc\_inode, destroy\_inode, read\_inode, write\_inode, sync\_fs, ...*

# VFS inode

- Uniquely identifies a file in a file system
- Access metadata (attributes) of the file (except name)

```
struct inode {  
    unsigned long i_ino;  
    umode_t i_mode;  
    uid_t i_uid;  
    gid_t i_gid;  
    kdev_t i_rdev;  
    loff_t i_size;  
    struct timespec i_atime;  
    struct timespec i_ctime;  
    struct timespec i_mtime;  
    struct super_block *i_sb;  
    struct inode_operations *i_op;  
    struct address_space *i_mapping;  
    struct list_head i_dentry;  
    ...  
}
```

*inode operations*



# VFS inode operations

Functions that operate on file & directory names and attributes

```
struct inode_operations {
    int (*create) (struct inode *, struct dentry *, int);
    struct dentry * (*lookup) (struct inode *, struct dentry *);
    int (*link) (struct dentry *, struct inode *, struct dentry *);
    int (*unlink) (struct inode *, struct dentry *);
    int (*symlink) (struct inode *, struct dentry *, const char *);
    int (*mkdir) (struct inode *, struct dentry *, int);
    int (*rmdir) (struct inode *, struct dentry *);
    int (*mknod) (struct inode *, struct dentry *, int, dev_t);
    int (*rename) (struct inode *, struct dentry *, struct inode *, struct dentry *);
    int (*readlink) (struct dentry *, char *,int);
    int (*follow_link) (struct dentry *, struct nameidata *);
    void (*truncate) (struct inode *);
    int (*permission) (struct inode *, int);
    int (*setattr) (struct dentry *, struct iattr *);
    int (*getattr) (struct vfsmount *mnt, struct dentry *, struct kstat *);
    int (*setxattr) (struct dentry *, const char *, const void *, size_t, int);
    ssize_t (*getxattr) (struct dentry *, const char *, void *, size_t);
    ssize_t (*listxattr) (struct dentry *, char *, size_t);
    int (*removexattr) (struct dentry *, const char *);
};
```

# VFS File operations

## Functions that operate on file & directory data

```
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char *, size_t, loff_t *);
    ssize_t (*aio_read) (struct kiocb *, char *, size_t, loff_t);
    ssize_t (*write) (struct file *, const char *, size_t, loff_t *);
    ssize_t (*aio_write) (struct kiocb *, const char *, size_t, loff_t);
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, int datasync);
    int (*aio_fsync) (struct kiocb *, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*lock) (struct file *, int, struct file_lock *);
    ssize_t (*readv) (struct file *, const struct iovec *, unsigned long, loff_t *);
    ssize_t (*writev) (struct file *, const struct iovec *, unsigned long, loff_t *);
    ssize_t (*sendfile) (struct file *, loff_t *, size_t, read_actor_t, void *);
    ssize_t (*sendpage) (struct file *, struct page *, int, size_t, loff_t *, int);
    unsigned long (*get_unmapped_area)(struct file *, unsigned long, unsigned long,
                                     unsigned long, unsigned long);
};
```

# VFS File operations

Not all functions need to be implemented!

Example: The same `file_operations` are used for a character device driver

```
struct file_operations mydriver_fops = {  
    .owner = MYFS_MODULE;  
    .open = myfs_open;          /* allocate resources */  
    .read = myfs_read_file;  
    .write = myfs_write_file;  
    .release = myfs_release;    /* release resources */  
    /* llseek, readdir, poll, mmap, readv, etc. not implemented */  
};
```

```
register_filesystem(&myfs_type)
```



```
static struct file_system_type myfs_type = {  
    .owner  = THIS_MODULE;  
    .name   = "myfs";  
    .get_sb = myfs_get_super,  
    .kill_sb = myfs_kill;  
};
```

# The End